

Claude GAP Press Release & FAQs

Press Release

Anthropic launches Governed Action Proxy, bringing autonomous Claude agents into production at regulated enterprises.

May 14, 2026 at 9:00 AM PT

Compliance teams at banks, hospitals, and law firms can now approve what an AI does before it does it. GAP intercepts high-risk Claude agent actions, runs them through customer-defined policy rules, and generates a cryptographic audit trail for every decision.

Available today for Claude Enterprise customers, with no code changes or agent retraining required, working with any tool Claude can call through Model Context Protocol, plugging in to the enterprise systems customers already have.

SAN FRANCISCO (BUSINESS WIRE) May 14, 2026. Anthropic today announced the general availability of Governed Action Proxy (GAP), a new compliance layer for Claude agents. GAP gives regulated enterprises the controls they need to deploy autonomous AI on workflows that touch real money, real patient data, and real legal exposure.

"This is the unlock that regulated industries have been waiting for," said Dario Amodei, CEO of Anthropic. "Our customers in finance, healthcare, and legal have never doubted that Claude could do the work. They doubted whether they could prove to a regulator that the work was done correctly. GAP makes that proof automatic."

For two years, enterprise leaders in finance, healthcare, and legal have faced the same wall. Their agents can draft a wire transfer, modify a clinical trial protocol, or update a customer record on their own. They cannot let those agents actually do it. A 2025 Gartner survey found that 40 percent of agentic AI projects are expected to be canceled by 2027, with unclear business value and governance gaps driving most cancellations. The deeper problem underneath those numbers is trust. Compliance teams cannot sign off on systems they cannot inspect.

Designed for the Workflows Auditors Actually Worry About

GAP intercepts the tool calls a customer flags as high risk. A wire transfer above a certain threshold. A change to a clinical trial protocol. A bulk write to a customer database. Each intercepted action runs through a policy engine the customer writes in plain text, the same way they write internal procedures today. Low-risk actions pass through. High-risk actions route to a named human approver, with the full reasoning chain attached so the reviewer can see what the agent was trying to do and why.

Every Action Logged, Every Decision Auditable

Every decision GAP makes lands in a cryptographic audit log. Approved actions, denied actions, escalations, the policy that fired, the human who signed off, and the agent's underlying reasoning. The audit trail is tamper-evident by design, so a customer can hand a year of activity to a regulator and prove that nothing has been altered after the fact. Compliance is logged as the work happens, not assembled after the fact.

Up and Running in a Day

Customers install GAP, point their existing MCP connections at it, and write their first policy the same day. The policy language sits on top of Open Policy Agent, an open-source standard already used inside most enterprise security teams. Most compliance and platform groups can author their first policy in a plain text file using patterns they already write today. No code changes. No new connectors. No agent retraining.

"Before GAP, our agents could draft a sanctions review but a person had to copy the recommendation into the actual system," said the Head of AI Governance at a top-five US bank. "Now the agents do the full review themselves and my compliance team can see every decision in real time. For the first time, I can hand a regulator a complete record of what our AI did and why, without spending a week assembling it."

"We've wanted to put Claude on protocol amendments for over a year, and our medical safety officers always said the same thing. Show me how I prove this to the FDA," said the VP of Clinical Operations at a Fortune 500 pharmaceutical company. "GAP solved that on day one. The policy gate routes every protocol change to the right safety officer, and the audit log captures the regulatory fields we need to file. Our team called it the most useful AI feature we've seen this year."

GAP is available today for Claude Enterprise customers. Customers can start a pilot by contacting their Anthropic account team or visiting anthropic.com/gap.

External FAQs

1. How do I start using GAP?

If you are an Anthropic Enterprise customer, contact your account team to enable GAP on your tenant. You will see a new admin section in the Anthropic console where you can connect GAP to your existing MCP servers and write your first policy.

2. What kinds of actions does GAP intercept?

You decide. GAP intercepts any tool call you flag as high-risk in your policy file. Common examples are wire transfers above a certain dollar amount, changes to patient records, modifications to clinical trial protocols, and bulk writes to production databases. Anything Claude can do through MCP, GAP can govern.

3. Do I have to write policies from scratch?

No. GAP ships with a starter policy library covering common patterns in finance, healthcare, and legal services. Most customers start with the library, customize a few rules to match their internal procedures, and build out from there over time.

4. Who gets the approval requests when an action is escalated?

You define the approval routing in your policy. You can route to a specific named individual, a role like Tier 2 Compliance, a Slack channel, or an email distribution list. For high-stakes actions you can require approval from more than one person.

5. How fast does GAP respond?

GAP adds minimal latency for actions that pass through automatically. For escalated actions, the wait time depends on how quickly your approvers respond, which is the same dynamic you have today with any human-in-the-loop workflow.

6. Can I export the audit log for regulators or internal auditors?

Yes. The audit log can be exported in standard formats including PDF, CSV, and SIEM-compatible JSON. The cryptographic signatures travel with the export so auditors can verify the log has not been altered.

7. Does GAP work with AI agents other than Claude?

GAP is purpose-built for Claude. The deepest functionality, including visibility into the agent's reasoning chain, depends on Claude's inference layer. Customers running mixed-AI deployments typically use GAP for Claude-based workflows and rely on their other vendors' governance tools for the rest.

8. How is my policy data and audit log secured?

GAP runs inside your Claude Enterprise tenant and follows the same security model as the rest of your Anthropic deployment. Your policies and audit logs are not shared with Anthropic, are not used to train models, and remain in the region you have configured for data residency.

Internal FAQs

1. What will customers like most about GAP?

The shift from advisory mode to autonomous mode on workflows their compliance team actually approves. Customers have been running Claude as a draft generator for two years. GAP lets them put Claude on the actual work without losing the audit trail their regulators require.

2. What will customers like least about GAP?

Writing the first policy. Customers in our research consistently said the policy language is clear once they see an example, but figuring out which actions to gate is harder than they expected. We are addressing this with a starter policy library that covers common patterns in each target industry, but the work of mapping a customer's specific procedures onto GAP rules will remain a real implementation lift.

3. Why are we launching this now?

Two things are converging. Gartner expects 40 percent of enterprise applications to include task-specific AI agents by end of 2026, up from under 5 percent last year. At the same time, the same research firm projects that 40 percent of agentic AI pilots will be canceled by 2027, largely because customers cannot get their compliance teams to sign off on production deployment. That gap between agent capability and agent governance is the opening GAP is built for.

4. Which of our strategic priorities does GAP support?

Owning the enterprise agentic layer. This is one of the two or three battles we have to win in the next phase of the company, alongside continued model quality investment and the Claude Code franchise. GAP is what turns Claude from the best model into the only agent a compliance team will sign off on.

5. What is our moat?

GAP sits on top of MCP, which we control. That gives us visibility into the agent's full reasoning chain in a way a third-party governance vendor cannot replicate. The more GAP is used, the deeper the policy library, the integration patterns, and the audit format become for customers, building hard-to-replicate assets, and a durable form of competitive advantage.

6. How does GAP compare to what Microsoft and Salesforce offer?

Microsoft Copilot's governance lives at the data layer through Purview. It can tell you what data Copilot accessed. It cannot tell you why the agent decided to take a specific action, because Copilot's reasoning chain is less inspectable than Claude's. Salesforce Agentforce has stronger workflow integration into Salesforce's own product surface but does not generalize to actions outside the Salesforce ecosystem. GAP is the only governance layer that works across any tool an enterprise has connected through their MCP servers, which our data suggests can be in the hundreds for large customers.

7. Who is our target customer?

Compliance officers and platform engineering leaders at Fortune 500 firms in regulated industries. The primary pull comes from compliance. Compliance officers are accountable to regulators for what their AI systems do, and they currently cannot get that accountability with autonomous agents. Platform engineering is the secondary buyer because they are the ones who actually deploy the proxy.

8. What is our GTM strategy?

GAP is positioned as an add-on to Claude Enterprise, sold through our existing field team rather than a separate sales org. We target one high-risk workflow, prove the value and functionality of GAP in the first quarter, then expand to others.

9. What does GAP depend on from the MCP roadmap?

GAP requires MCP servers to surface tool calls in a structured format before execution. Most connectors in our marketplace already meet this bar. The remaining ones need a small protocol update that the MCP working group has already scoped. If that update slips past our launch date, GAP ships with the connectors that are ready and will support the remaining servers in a follow-up release.

10. What happens when GAP itself fails?

Two failure modes. GAP approves something it should have blocked, or GAP blocks something it should have approved. For the first, the audit log captures exactly which policy fired so the customer can update the rule. Anthropic is not liable for the underlying action, which is consistent with competitive compliance tools. For the second, the customer gets a clear escalation path with manual override, which is also logged. We track false positives as a quality bug, and use corrected signals to improve the reasoning of new models.

11. How do we measure success?

The North Star is Core Workflow Penetration Rate, the percentage of a customer's mission-critical processes running through GAP-governed Claude agents. Target is 25 percent within 12 months of deployment. We also watch whether customers actually write their own policies (a sign they trust the system enough to invest in it) and how often compliance teams query the audit log (a sign the audit log is providing real value). Our guardrail metric is latency added by GAP, which we cap aggressively so we do not slow down the agent.

12. What are the top risks?

A high-profile customer failure that ends in regulatory action. We mitigate this through a phased rollout where early customers get hands-on deployment support. Second risk is Microsoft bundling equivalent governance into Copilot at zero marginal cost. Our defense is the integration depth into Claude's reasoning, which a Copilot equivalent cannot match. Third is EU AI Act enforcement guidance changing in a way that breaks our policy framework. The open-source policy layer lets customers adapt rules without waiting on us.

13. How will competitors respond?

Microsoft will respond first, likely through Purview. Their advantage is bundling. Their disadvantage is that Copilot's reasoning is harder to inspect than Claude's. Google's response will lag because Vertex AI's governance story is still focused on data residency rather than agent action governance. OpenAI is the wild card. Their enterprise team has been investing in compliance tooling, and an equivalent product within a year is plausible.

14. Why not let customers build this themselves?

We expected some will. The customers who can build their own governance layer (large banks with internal AI infrastructure teams) are a small minority of the regulated segment. Most customers want a buy decision, not a build decision. For the build-it-themselves group, GAP's policy library and audit format become reference implementations they can pull from, which is part of why the policy layer is open source.

15. Why does the prototype not include a deny option for escalated transfers?

The v1 prototype focuses on demonstrating the first half of the governance loop: detecting a high-risk transfer, escalating it for human review, and logging that an approval step is required. However, a production version of GAP would need a clear deny path for approvers. Without denial, the product cannot fully support compliance review because reviewers need the ability to block actions that violate policy, appear suspicious, or lack sufficient context.

In production, denied actions would be written to the audit log with the policy that triggered review, the approver's identity, the denial reason, timestamp, and any follow-up action required. This matters because regulators and internal auditors care not only about approved actions, but also about which actions were stopped and why.

Appendix: Known Limitations

This section documents the explicit trade-offs and scope boundaries of the GAP v1 prototype. These are not bugs. They are decisions we made about what to build now and what to defer.

1. The policy engine is a single hardcoded rule.

The prototype intercepts wire transfers and auto-approves anything under \$10,000. That is the entire policy logic. A production GAP would replace this with Open Policy Agent and a customer-authored policy file supporting arbitrary rules across any tool call type. We chose the hardcoded threshold to keep the prototype focused on demonstrating the escalation pattern (intercept, route, log) rather than the policy authoring experience.

2. The chat is simulated, not connected to a real model.

The prototype runs canned Claude responses on a 250ms delay rather than calling the Anthropic API. This let us build and demo the interception flow without burning tokens or dealing with non-determinism during development. A production GAP would sit between real Claude inference and real MCP tool calls.

3. The tools GAP intercepts are stubbed.

The only tool call the prototype handles is a simulated `execute_wire_transfer`. There are no real MCP server connections, no actual external systems, and no real money or data at risk. A production GAP would intercept any MCP tool call from any connected server.

4. The audit log is in-memory and not cryptographic.

The Agent Pane logs actions for the current session only. Closing the browser clears the log. The press release describes a cryptographic, tamper-evident audit trail. That is a production requirement we did not build into the prototype because it adds storage infrastructure that does not change the user-facing demo.

5. The approver is a single static role.

When an action escalates, the prototype shows the request in the Action Pane with no logic about who specifically should approve, multi-party approval, escalation timeouts, or notification routing. A production GAP would handle approver lookup, Slack/email integration, and timeout fallback rules.

6. There is no policy authoring UI.

A real customer would write policies in a text editor and version them in their existing tooling. The prototype skips this entirely. We are not making claims about the authoring experience until we have done user research with compliance and platform teams.

7. No multi-tenancy, no auth, no admin controls.

Anyone can load the prototype and trigger any flow. There is no concept of users, organizations, roles, or permissions. A production GAP would inherit the tenant model from Claude Enterprise.

8. The prototype demonstrates the pattern, not the integration.

The MCP-layer integration is the most important architectural claim in the press release and the most expensive thing to build. We deferred it because the user-facing demo is the same with or without it. A real implementation requires changes to how MCP servers surface tool calls before execution, which the press release acknowledges as a dependency on the MCP roadmap.

What we are intentionally not building in v1, and why

1. Open Policy Agent integration. Building a production policy engine before validating that compliance teams will use a plain-text rule format is premature. Our next research milestone is showing the current prototype to three compliance officers and asking them to author a policy aloud.

2. A cryptographic audit format. The audit log primitive is easy to add once we know what fields regulators actually want logged. That research happens with auditors and external counsel, not in our prototype.

3. Approval workflow integrations with Slack, email, ServiceNow. These are channel integrations, not core systems. We will pick the first one based on what our pilot customers already use.

4. Pricing and packaging. We are not committed to a specific price point, with pricing finalized after the pilot program.